

E-commerce Application Optimization Report

1. Identified Bottlenecks

Our analysis of the initial E-commerce application revealed several performance bottlenecks:

- Unnecessary delay in Order total calculation:** The `calculateTotal()` method in the `Order` class included an artificial delay, significantly slowing down order processing.
- Inefficient product loading:** The `Inventory` class loaded a large number of products into memory at startup, causing high memory usage and slow initialization.
- Inefficient product search:** The `searchProducts()` method in the `Inventory` class used a linear search algorithm, resulting in slow search operations, especially with a large number of products.
- Slow product retrieval by ID:** The `getProductById()` method in the `Inventory` class used a linear search, leading to poor performance for individual product lookups.
- Sequential order processing:** The `OrderProcessor` class processed orders sequentially, not utilizing available system resources effectively.

2. Performance Improvements

To address these bottlenecks, we implemented the following improvements:

1. Optimized Order total calculation:

- Removed the artificial delay.
- Utilized Java streams for efficient calculation.

```
public double calculateTotal() {  
    return products.stream().mapToDouble(Product::getPrice).sum();  
}
```

2. Implemented lazy loading for products:

- Products are now loaded in batches as needed, reducing memory usage and improving startup time.

```
private void loadProductsBatch(int start, int batchSize) {  
    for (int i = start; i < start + batchSize; i++) {  
        Product product = new Product("P" + i, "Product " + i, Math.random()  
            * 100, "Description " + i);  
        productsMap.put(product.getId(), product);  
        updateSearchIndex(product);  
    }  
}
```

3. Improved product search with indexing:

- Implemented an in-memory search index for faster product searches.

```
public List<Product> searchProducts(String keyword) {
    return searchIndex.getDefault(keyword.toLowerCase(),
Collections.emptyList());
}
```

4. Optimized product retrieval:

- Used a `ConcurrentHashMap` to store products with their IDs as keys, allowing O(1) lookup time.

```
public Product getProductById(String id) {
    return productsMap.get(id);
}
```

5. Implemented parallel order processing:

- Utilized `CompletableFuture` for asynchronous and parallel order processing.
- Optimized I/O operations with `BufferedWriter`.

```
public void processOrders() {
    List<CompletableFuture<Void>> futures = orders.stream()
        .filter(order -> order.getStatus().equals("NEW"))
        .map(this::processOrderAsync)
        .collect(Collectors.toList());

    CompletableFuture.allOf(futures.toArray(new
CompletableFuture[0])).join();
}
```

3. Performance Comparison

While we don't have exact metrics due to the nature of this lab, we can estimate the performance improvements:

| Operation | Before Optimization | After Optimization | Estimated Improvement |
|-------------------------|-------------------------|--------------------|---------------------------------------|
| Order Total Calculation | O(n) + artificial delay | O(n) | ~100x faster |
| Product Loading | O(n) at startup | O(1) lazy loading | Instant startup, reduced memory usage |

| Operation | Before Optimization | After Optimization | Estimated Improvement |
|-------------------------|---------------------|---|--|
| Product Search | $O(n)$ | $O(1)$ average case | ~1000x faster for large inventories |
| Product Retrieval by ID | $O(n)$ | $O(1)$ | ~1000x faster for large inventories |
| Order Processing | Sequential, $O(n)$ | Parallel, $O(n/k)$ where k is number of cores | 2-8x faster depending on available cores |

Note: Actual performance gains may vary based on the specific hardware and dataset sizes.

4. Adherence to 12-Factor Principles

Let's evaluate our application against the 12-factor app methodology:

- Codebase:** ☒ The application can be version-controlled in a single repository.
- Dependencies:** ☒ No explicit dependency management (e.g., no Maven or Gradle). Improvement: Implement a build tool like Maven or Gradle to manage dependencies.
- Config:** ☒ Configuration is hardcoded. Improvement: Use environment variables or config files for settings like file paths and thread pool sizes.
- Backing Services:** ☒ Not applicable in the current version, but could be improved by treating the file system as an attached resource.
- Build, Release, Run:** ☒ No clear separation of build and run stages. Improvement: Implement a build process and create distinct release packages.
- Processes:** ☒ The application is stateless and shares nothing between runs.
- Port Binding:** ☒ Not applicable (not a web service).
- Concurrency:** ☒ The application can scale out via the process model with the `ExecutorService`.
- Disposability:** ☒ No fast startup or graceful shutdown implemented. Improvement: Implement proper shutdown hooks and startup optimizations.
- Dev/Prod Parity:** ☒ No clear separation of development and production environments. Improvement: Use configuration to differentiate between environments.
- Logs:** ☒ Logs are not treated as event streams. Improvement: Implement a logging framework and treat logs as streams of events.
- Admin Processes:** ☒ Not applicable in the current version.

Conclusion

Our optimization efforts have significantly improved the performance of the E-commerce application, particularly in areas of product management, search, and order processing. The application now scales better

and can handle larger datasets more efficiently.

However, there's still room for improvement, especially in adhering to all 12-factor app principles. Future work should focus on improving dependency management, configuration handling, build processes, and logging to make the application more robust and maintainable.